

# Nuxt + Cloudflare Workers + Hyperdrive + Prisma

## Architecture, Setup, and Integration Guide

### 1. Overview & Goals

This document describes how to build and integrate a backend API for a marketplace-style Nuxt 3 application using Cloudflare Workers, Cloudflare Hyperdrive, Amazon RDS (PostgreSQL), and Prisma. The API is split into two logical surfaces, both implemented in a single Worker script:

- **Private API:** Consumed only by your own Nuxt frontend. Authenticated via session JWT stored in an HttpOnly cookie, and exposed on the hostname `api.example.com`.
- **Public API:** Exposed to third-party integrators. Authenticated via per-account API keys passed in an `Authorization: ApiKey ...` header, and exposed on `public-api.example.com`.

Backend data is stored in a PostgreSQL database hosted on Amazon RDS. Cloudflare Hyperdrive provides a pooled, serverless-friendly connection between Cloudflare Workers and the RDS instance. Prisma is used as the ORM/data access layer, with its driver adapters feature to run inside Workers.

#### ***High-Level Component Diagram (conceptual)***

- Nuxt app (`app.example.com`) → calls private API on `api.example.com` using cookies (session JWT).
- Third-party clients → call public API on `public-api.example.com` using API keys.
- Both hostnames map to the same Cloudflare Worker script.
- Worker → connects to Amazon RDS Postgres through Cloudflare Hyperdrive using Prisma.

## 2. Infrastructure Setup

### 2.1 PostgreSQL on Amazon RDS

- 1) In AWS, create a PostgreSQL RDS instance (choose appropriate size and storage for development or production).
- 2) Configure security groups / networking so that only Cloudflare / Hyperdrive can reach the database. For early development you may temporarily allow your IP, but plan to tighten this for production.
- 3) Create a dedicated database user for your application (do not use the master user in production).
- 4) Note the connection string in the form: postgresql://USER:PASSWORD@HOST:PORT/DB\_NAME. This is used by Prisma migrations and by Hyperdrive.

### 2.2 Cloudflare Worker Project (Wrangler)

```
# Create a new Workers project in TypeScript
npm create cloudflare@latest marketplace-api-worker -- --type=hello-world --ts=true --git=true --deploy=false
cd marketplace-api-worker
```

This creates a minimal Worker project. You will extend it with Prisma, routing, and authentication logic.

### 2.3 Enable Node.js Compatibility and Hyperdrive

Edit wrangler.toml to enable Node.js compatibility (needed for Prisma and the pg driver) and to bind your Hyperdrive configuration.

```
# wrangler.toml (core parts only)
name = "marketplace-api-worker"
main = "src/index.ts"
compatibility_date = "2024-09-23"
compatibility_flags = ["nodejs_compat"]

[[hyperdrive]]
binding = "HYPERDRIVE"
id = "<your-hyperdrive-id-here>"
```

The binding value (HYPERDRIVE) becomes env.HYPERDRIVE inside your Worker code. The id is obtained when you create a Hyperdrive config in the Cloudflare dashboard.

### 2.4 Hostnames and Routes

Configure two hostnames in Cloudflare DNS and route them to the same Worker:

```
[env.production]
routes = [
  { pattern = "api.example.com/*", zone_name = "example.com" },
  { pattern = "public-api.example.com/*", zone_name = "example.com" }
]
```

Both hostnames execute the same Worker script. The script will inspect url.hostname to decide whether to handle the request as private (Nuxt) or public (third-party API).

## 3. Prisma Setup (Workers + Hyperdrive)

### 3.1 Install Prisma and Dependencies

```
npm install -D prisma
npx prisma init

npm install pg@>8.13.0 @prisma/adapter-pg
npm install -D @types/pg
```

This initializes a Prisma project and installs the PostgreSQL driver plus the Prisma driver adapter for pg. Prisma will be used both locally (for migrations) and in the Worker (via the adapter).

### 3.2 Prisma Schema (Marketplace Models)

Edit `prisma/schema.prisma` to define the generator, datasource, and basic marketplace models. The schema below is intentionally simplified and can be extended as needed.

```
generator client {
  provider      = "prisma-client-js"
  previewFeatures = ["driverAdapters"]
}

datasource db {
  provider = "postgresql"
  url      = env("DATABASE_URL")
}

model User {
  id        String    @id @default(cuid())
  email     String    @unique
  password  String
  role      String    // "buyer" | "seller" | "admin"
  createdAt DateTime  @default(now())
  orders    Order[]
  apiKeys   ApiKey[]
}

model Product {
  id        String    @id @default(cuid())
  title     String
  description String
  priceCents Int
  stock     Int
  createdAt DateTime @default(now())
  updatedAt DateTime @updatedAt
  orderItems OrderItem[]
}

model Order {
  id        String    @id @default(cuid())
  userId    String
  user      User       @relation(fields: [userId], references: [id])
  status    String    // "PENDING" | "PAID" | "CANCELLED"
  createdAt DateTime  @default(now())
  items     OrderItem[]
}

model OrderItem {
  id        String    @id @default(cuid())
  orderId   String
  productId String
  quantity  Int
  priceCents Int

  order    Order    @relation(fields: [orderId], references: [id])
  product  Product  @relation(fields: [productId], references: [id])
}

model ApiKey {
```

```
id      String    @id @default(cuid())
userId String
user    User      @relation(fields: [userId], references: [id])
name    String
keyHash String    @unique
scopes  String[]
rateLimit Int      @default(10000)
createdAt DateTime @default(now())
lastUsedAt DateTime?
disabledAt DateTime?

}
```

### 3.3 Local .env and Migrations

Prisma uses a local .env file to find DATABASE\_URL for migrations and generation. This value should be a direct PostgreSQL connection string to your RDS instance (or to a local Postgres during early development).

```
# .env (local dev / CI for migrations)
DATABASE_URL="postgresql://USER:PASSWORD@HOST:5432/DB_NAME"
```

Add scripts to package.json:

```
"scripts": {
  "prisma:migrate": "prisma migrate dev",
  "prisma:generate": "prisma generate --no-engine"
}
```

Run the migrations and generate the client:

```
npm run prisma:migrate
npm run prisma:generate
```

## 4. Worker Code Structure

### 4.1 Suggested Directory Layout

```
src/
  index.ts          # Worker entrypoint: routing & dispatch
  prismaClient.ts  # Prisma client helper using Hyperdrive adapter

  authPrivate.ts    # Cookie-based session auth (Nuxt)
  authPublic.ts    # API key auth (third parties)

  routes/
    private/
      me.ts          # GET /internal/me
      products.ts    # GET /internal/products, etc.
      orders.ts      # POST /internal/orders, etc.
      auth.ts        # POST /internal/auth/login, logout

    public/
      products.ts    # GET /public/v1/products
      orders.ts      # e.g. GET /public/v1/orders (scoped)
```

### 4.2 Prisma Client Helper (Workers + Hyperdrive)

The helper below creates a PrismaClient instance using the Hyperdrive binding. For simplicity, it creates a new client per request. Hyperdrive handles pooling at the network edge.

```
// src/prismaClient.ts
import { PrismaClient } from "@prisma/client"
import { PrismaPg } from "@prisma/adapters"

export interface Env {
  HYPERDRIVE: { connectionString: string }
  JWT_SECRET: string
}

export function createPrisma(env: Env): PrismaClient {
  const adapter = new PrismaPg({
    connectionString: env.HYPERDRIVE.connectionString,
  })

  return new PrismaClient({ adapter })
}
```

### 4.3 Worker Entrypoint (Hostname-Based Routing)

```
// src/index.ts
import type { Env } from "./prismaClient"
import { authenticatePrivate } from "./authPrivate"
import { authenticatePublicKey } from "./authPublic"
import { handlePrivateApi } from "./routes/private/index"
import { handlePublicApi } from "./routes/public/index"

function json(body: unknown, status = 200): Response {
  return new Response(JSON.stringify(body), {
    status,
    headers: { "Content-Type": "application/json" },
  })
}

export default {
  async fetch(request: Request, env: Env, ctx: ExecutionContext): Promise<Response> {
    const url = new URL(request.url)
    const hostname = url.hostname

    if (hostname === "api.example.com") {
      const session = await authenticatePrivate(request, env)
      if (!session) return json({ error: "Unauthorized" }, 401)
    }
  }
}
```

```
    return handlePrivateApi(request, env, ctx, session)
}

if (hostname === "public-api.example.com") {
  const apiContext = await authenticatePublicKey(request, env)
  if (!apiContext) return json({ error: "Invalid API key" }, 401)
  return handlePublicApi(request, env, ctx, apiContext)
}

return json({ error: "Unknown host" }, 404)
},
} satisfies ExportedHandler<Env>
```

## 5. Authentication Flows

### 5.1 Private API (Nuxt) – Session JWT in Cookie

The private API is used only by your own Nuxt app. Authentication is via a JWT stored in an HttpOnly cookie named session on the api.example.com hostname (or .example.com domain).

Login flow:

- 1) User submits email/password from the Nuxt app to POST /internal/auth/login on api.example.com.
- 2) Worker verifies credentials with Prisma.
- 3) Worker signs a JWT with sub = user.id and role = user.role using env.JWT\_SECRET.
- 4) Worker returns a response with Set-Cookie: session=...; HttpOnly; Secure; SameSite=Lax; Domain=.example.com.
- 5) Subsequent requests from Nuxt include this cookie and are authenticated by authenticatePrivate.

```
// src/authPrivate.ts
import * as jose from "jose"
import type { Env } from "./prismaClient"

export type Session = { userId: string; role: string }

export async function authenticatePrivate(request: Request, env: Env): Promise<Session | null> {
  const cookieHeader = request.headers.get("Cookie") || ""
  const match = cookieHeader.match(/session=([^;]+)/)
  if (!match) return null

  const token = decodeURIComponent(match[1])

  try {
    const { payload } = await jose.jwtVerify(
      token,
      new TextEncoder().encode(env.JWT_SECRET),
    )
    return {
      userId: String(payload.sub),
      role: String(payload.role ?? "user"),
    }
  } catch {
    return null
  }
}
```

### 5.2 Public API – API Keys

The public API is accessed by third-party clients via API keys. Keys are created and managed via your internal dashboard (using the private API). Only a hash of each key is stored in the database to avoid leaking secrets.

- Request header format: Authorization: ApiKey <key>.
- Worker hashes the presented key and looks it up in the ApiKey table.
- If found and not disabled, the request is associated with the owning user/account and the scopes stored on the record (for example, ["read\_products"]).

```
// src/authPublic.ts
import { createPrisma, Env } from "./prismaClient"

export type ApiContext = { apiKeyId: string; userId: string; scopes: string[] }

export async function hashApiKey(raw: string): Promise<string> {
  // Example: HMAC-SHA256 with a server-side secret, bcrypt, or argon2.
  // Implementation detail left to the developer.
  return raw // placeholder
}

export async function authenticatePublicApiKey(
  request: Request,
  env: Env,
): Promise<ApiContext | null> {
```

```
const auth = request.headers.get("Authorization") || ""
const m = auth.match(/^ApiKey\s+(.+)$/i)
if (!m) return null

const apiKey = m[1].trim()
if (!apiKey) return null

const keyHash = await hashApiKey(apiKey)

const prisma = createPrisma(env)
const record = await prisma.apiKey.findUnique({
  where: { keyHash },
})

if (!record || record.disabledAt) return null

return {
  apiKeyId: record.id,
  userId: record.userId,
  scopes: record.scopes,
}
}
```

## 6. Example Route Handlers

### 6.1 Private: List Products

The private products endpoint is used by the Nuxt app to display products in the internal UI (for example, with admin-only fields or draft listings).

```
// src/routes/private/products.ts
import { createPrisma, Env } from "../../prismaClient"
import type { Session } from "../../authPrivate"

function json(body: unknown, status = 200): Response {
  return new Response(JSON.stringify(body), {
    status,
    headers: { "Content-Type": "application/json" },
  })
}

export async function listPrivateProducts(
  request: Request,
  env: Env,
  session: Session,
): Promise<Response> {
  const url = new URL(request.url)
  const page = Number(url.searchParams.get("page") || "1")
  const pageSize = Number(url.searchParams.get("pageSize") || "20")

  const prisma = createPrisma(env)

  const [items, total] = await Promise.all([
    prisma.product.findMany({
      skip: (page - 1) * pageSize,
      take: pageSize,
      orderBy: { createdAt: "desc" },
    }),
    prisma.product.count(),
  ])

  return json({ data: items, total, page, pageSize })
}
```

### 6.2 Public: List Products

The public products endpoint returns only the fields meant for third-party consumption and enforces whatever "public visibility" rules you decide (for example, only published products).

```
// src/routes/public/products.ts
import { createPrisma, Env } from "../../prismaClient"
import type { ApiContext } from "../../authPublic"

function json(body: unknown, status = 200): Response {
  return new Response(JSON.stringify(body), {
    status,
    headers: { "Content-Type": "application/json" },
  })
}

export async function listPublicProducts(
  request: Request,
  env: Env,
  api: ApiContext,
): Promise<Response> {
  const url = new URL(request.url)
  const page = Number(url.searchParams.get("page") || "1")
  const pageSize = Number(url.searchParams.get("pageSize") || "20")

  const prisma = createPrisma(env)

  const [items, total] = await Promise.all([
    prisma.product.findMany({
```

```
    skip: (page - 1) * pageSize,
    take: pageSize,
    orderBy: { createdAt: "desc" },
    where: {
      // add public visibility filters here (e.g. published = true)
    },
    select: {
      id: true,
      title: true,
      description: true,
      priceCents: true,
      // omit any internal-only fields
    },
  ),
  prisma.product.count({
    where: {
      // same filter as above
    },
  }),
]
)

return json({ data: items, total, page, pageSize })
}
```

## 7. Nuxt Integration

### 7.1 Runtime Config and Environment

Nuxt needs to know where to call the private API. This is done via runtime config. In production, NUXT\_PUBLIC\_API\_BASE\_URL should be set to https://api.example.com.

```
// nuxt.config.ts
export default defineNuxtConfig({
  runtimeConfig: {
    public: {
      apiUrl: process.env.NUXT_PUBLIC_API_BASE_URL || "https://api.example.com",
    },
  },
})
```

### 7.2 Composable for Private API Calls

This composable wraps \$fetch and always sends cookies (credentials: "include") so that the Worker can authenticate the user via the session cookie.

```
// composables/useApi.ts
export const useApi = () => {
  const config = useRuntimeConfig()
  const baseURL = config.public.apiUrl

  const get = <T>(url: string, options: any = {}) =>
    $fetch<T>(url, {
      baseURL,
      credentials: "include",
      ...options,
    })

  const post = <T>(url: string, body: any, options: any = {}) =>
    $fetch<T>(url, {
      method: "POST",
      baseURL,
      body,
      credentials: "include",
      ...options,
    })

  return { get, post }
}
```

### 7.3 Nuxt Page Example: Products List

```
// pages/products.vue
<script setup lang="ts">
const { get } = useApi()

const { data: products, error } = await useAsyncData("products", () =>
  get("/internal/products", {
    query: { page: 1, pageSize: 20 },
  }),
)
</script>

<template>
  <div>
    <h1>Products</h1>
    <div v-if="error">Failed to load products</div>
    <ul v-else-if="products">
      <li v-for="p in products.data" :key="p.id">
        {{ p.title }} - {{ p.priceCents / 100 }} USD
      </li>
    </ul>
  </div>
</template>
```

```
</template>
```

## **7.4 Nuxt Login Flow (High-Level)**

- 1) User fills in email/password in a Nuxt form.
- 2) Nuxt posts to POST /internal/auth/login on api.example.com using useApi().post.
- 3) Worker verifies credentials, sets the session cookie, and returns a 200 OK.
- 4) On success, Nuxt navigates the user to a protected page. Subsequent calls automatically include the cookie.

## 8. Local Development & Deployment

### 8.1 Local Development Workflow

- 1) Configure DATABASE\_URL in .env to point at a development RDS instance (or local Postgres).
- 2) Run npm run prisma:migrate to apply schema changes.
- 3) Run npm run prisma:generate to update the client.
- 4) Start the Worker locally with npx wrangler dev.
- 5) In the Nuxt project, set NUXT\_PUBLIC\_API\_BASE\_URL to the local wrangler dev URL (or use host rewrites).
- 6) Run npm run dev in the Nuxt project and develop against the Worker.

### 8.2 Deployment Checklist

Database:

- Production RDS instance created and secured.
- Prisma migrations run against production (manually or via CI).

Cloudflare:

- Hyperdrive configured and bound as HYPERDRIVE.
- Worker deployed with compatible compatibility\_date and nodejs\_compat flag.
- Routes configured for api.example.com and public-api.example.com.
- Secrets (JWT\_SECRET, any hashing secrets) configured as Worker environment variables.

Nuxt:

- NUXT\_PUBLIC\_API\_BASE\_URL set to https://api.example.com in production env.
- Nuxt deployed (for example, to Vercel, Netlify, or Cloudflare Pages).

Security & Observability:

- Rate limiting and logging implemented for both private and public APIs.
- Access logs and metrics monitored for anomalies.
- Database performance (connections, slow queries, storage) monitored in RDS.

This guide should give a developer enough information to build and operate the API service using Cloudflare Workers, Hyperdrive, and Prisma, and to integrate it cleanly with a Nuxt 3 frontend and a public third-party API surface.